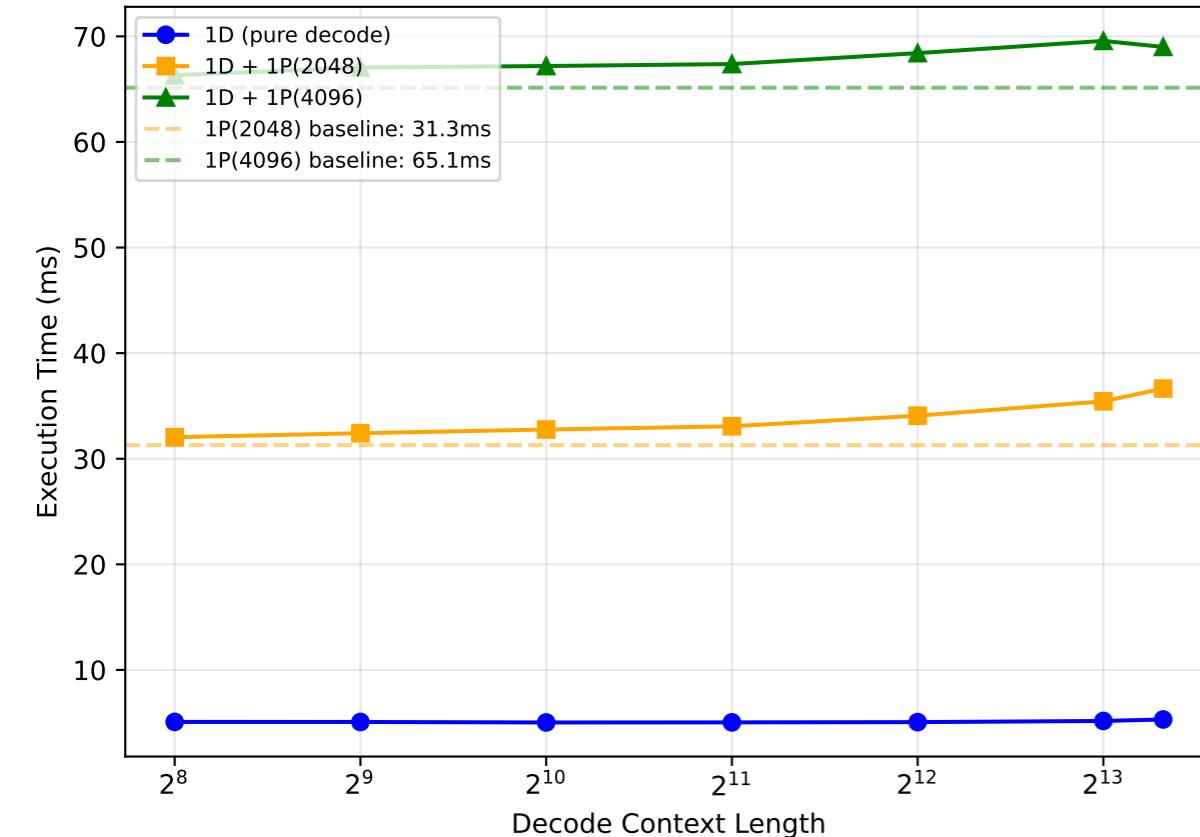
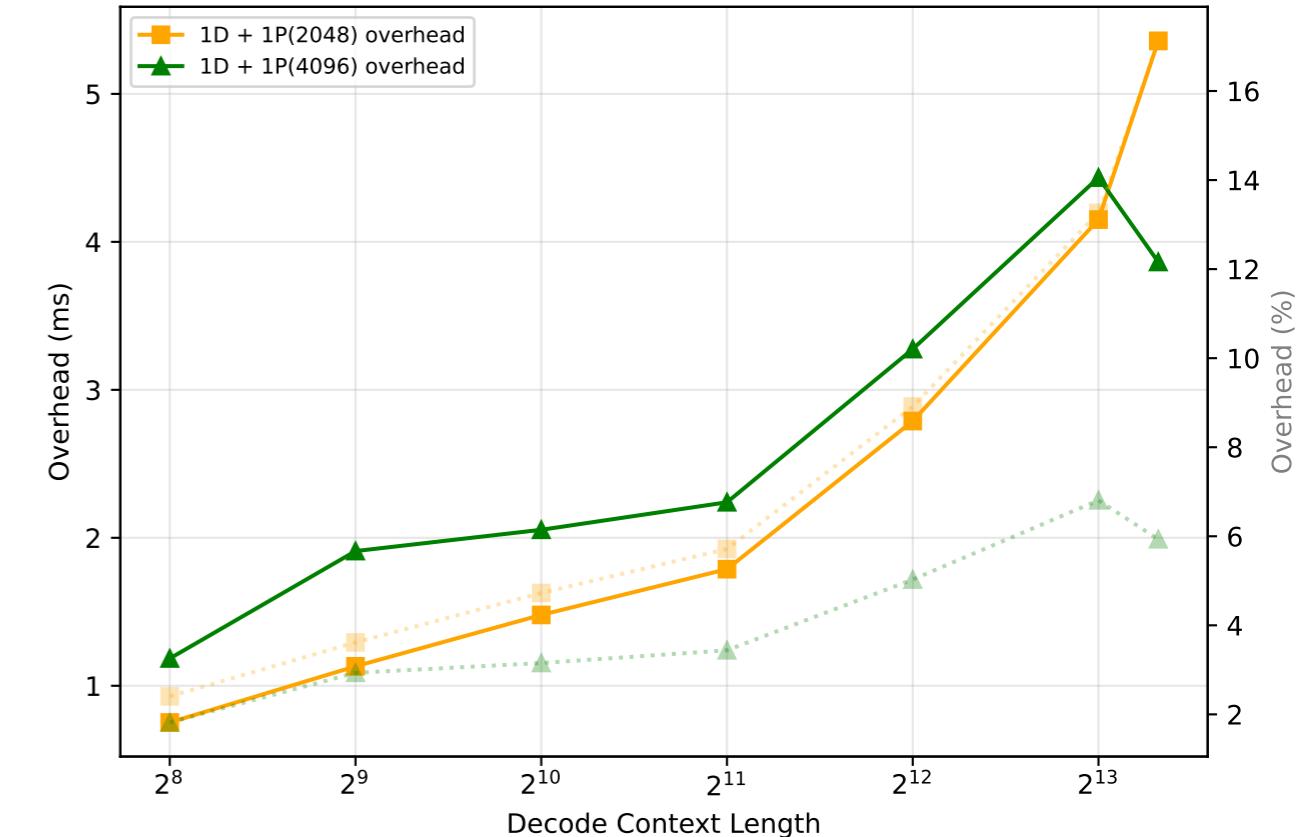


# vLLM Benchmark Results - Qwen/Qwen3-4B

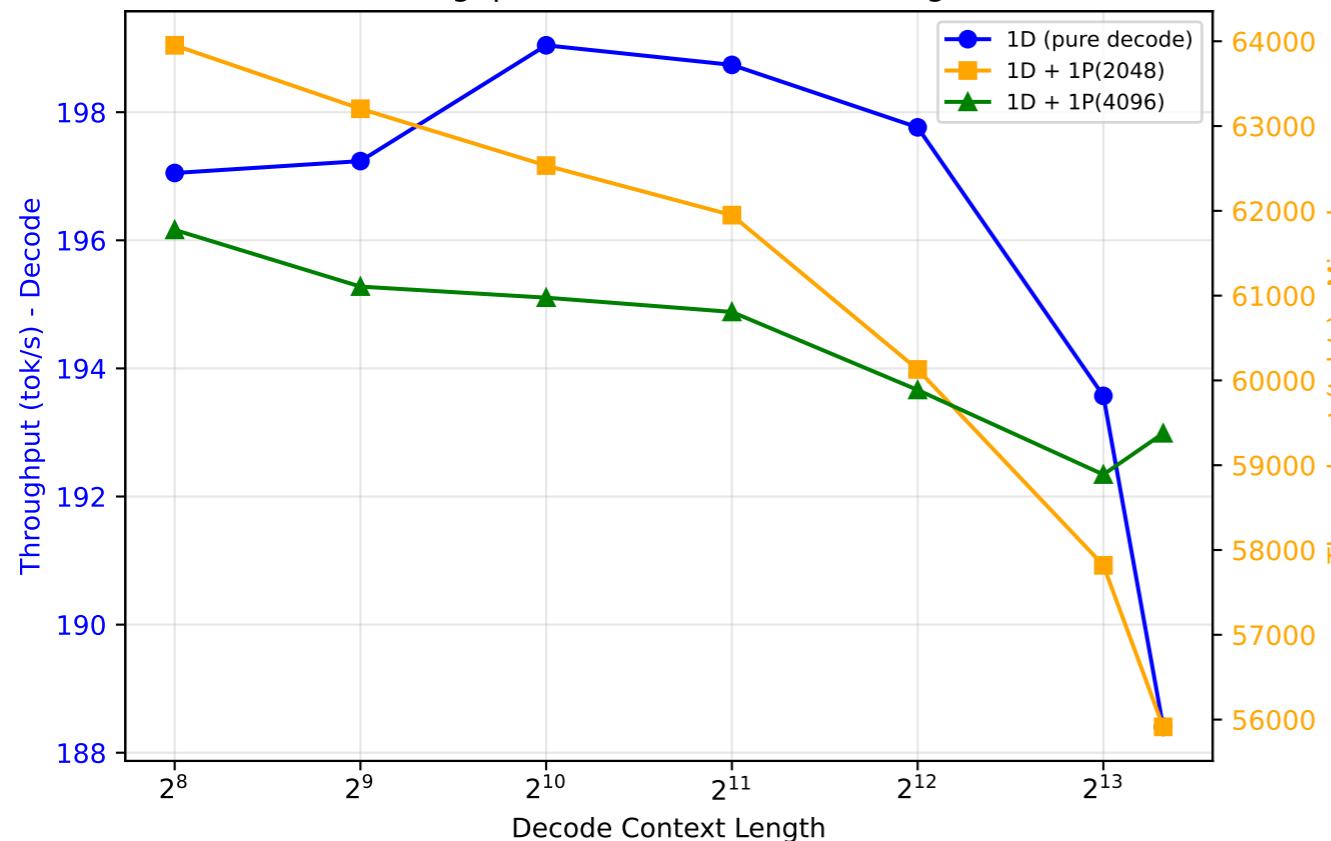
Execution Time vs Decode Context Length



Overhead of Adding 1 Decode to Prefill



Throughput vs Decode Context Length



Execution Time Comparison (Bar Chart)

